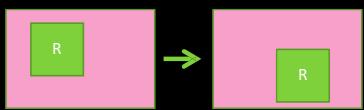


Special Projective Transformations

- Translation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Preserves:
 - Lengths/Areas
 - Angles
 - Orientation
 - Lines



Projective Transformations

- *Projective* transformations: for 2D images it's a 3x3 matrix applied to homogenous coordinates

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Special Projective Transformations

- Euclidean (Rigid body)

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & t_x \\ \sin(\theta) & \cos(\theta) & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Preserves:
 - Lengths/Areas
 - Angles
 - Lines



Special Projective Transformations

- Similarity (trans, rot, scale) transform

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a \cos(\theta) & -a \sin(\theta) & t_x \\ a \sin(\theta) & a \cos(\theta) & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Preserves:

- Ratios of Areas
- Angles
- Lines



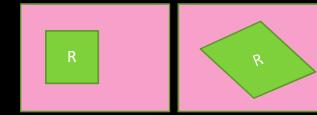
Special Projective Transformations

- Affine transform

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Preserves:

- Parallel lines
- Ratio of Areas
- Lines



Projective Transformations

- Remember, these are homogeneous coordinates

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} \cong \begin{bmatrix} sx' \\ sy' \\ s \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Projective Transformations

- General projective transform (or Homography)

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} \cong \begin{bmatrix} wx' \\ wy' \\ w \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Preserves:

- Lines
- Also cross ratios
(maybe later)



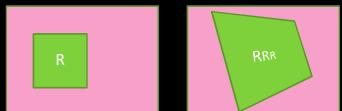
Projective Transformations

- General projective transform (or Homography)

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} \equiv \begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Preserves:

- Lines



- Also cross ratios
(maybe later)

Quiz 1

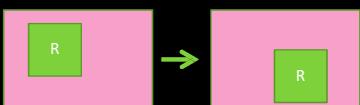
Suppose I told you the transform from image A to image B is a **translation**. How many pairs of corresponding points would you need to know to compute the transformation?

- a) 3
- b) 1
- c) 2
- d) 4

Quiz 1 – answer

- Translation: a 1 point transformation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$



Quiz 2

Suppose I told you the transform from image A to image B is **affine**. How many pairs of corresponding points would you need to know to compute the transformation?

- a) 3
- b) 1
- c) 2
- d) 4

Quiz 2 – answer

- Affine transform: a 3 point transformation
 - 6 unknowns – each point pair gives two equations

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

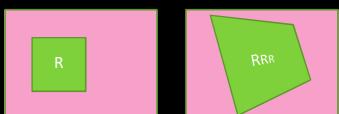
Quiz 3

Suppose I told you the transform from image A to image B is a **homography**. How many pairs of corresponding points would you need to know to compute the transformation?

- a) 3
- b) 1
- c) 2
- d) 4

Quiz 3 – answer

- Homography:
4 points



$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} \approx \begin{bmatrix} w'x' \\ w'y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Ain 431
Introduction to Computer Vision

Homographies and mosaics

Projective Transformations

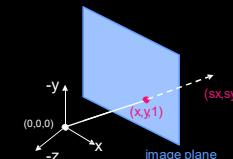
Projective transformations: for 2D images it's a 3x3 matrix applied to homogenous coordinates

$$\begin{bmatrix} w^* & x \\ w^* & y \\ w^* \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

The projective plane

What is the geometric intuition of using homogenous coordinates?

- A point in the image is a ray in projective space



The projective plane

Each *point* (x,y) on the plane (at $z=1$) is represented by a *ray* (sx,sy,s)

All points on the ray are equivalent:
 $(x, y, 1) \cong (sx, sy, s)$

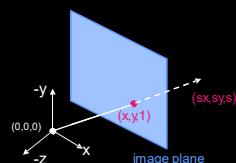
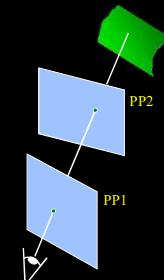


Image reprojection

Basic question:

How to relate two images from the same camera center?

How to map a pixel from projective plane PP1 to PP2?



Source: Alysha Efros

Image reprojection

Answer

- Cast a ray through each pixel in PP1
- Draw the pixel where that ray intersects PP2

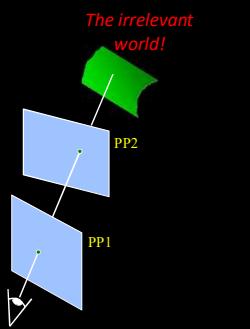
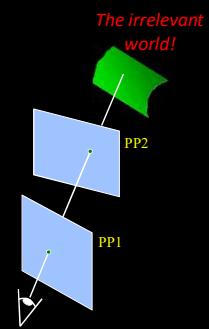


Image reprojection

Observation:

- Rather than thinking of this as a 3D reprojection, think of it as a 2D **image warp** from one image (plane) to another (plane).



Application: Simple mosaics

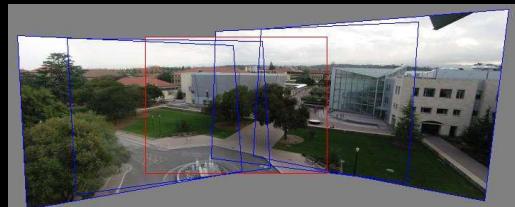


Image from http://graphics.cs.cmu.edu/courses/15-463/2010_fall/

How to stitch together a panorama (a.k.a. mosaic)?

Basic Procedure

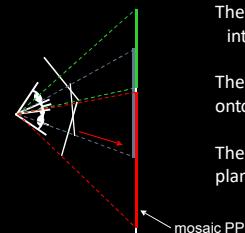
- Take a sequence of images from the same position
 - > Rotate the camera about its optical center
- Compute transformation between second image and first
- Transform the second image to overlap with the first
- Blend the two together to create a mosaic
- (If there are more images, repeat)

But wait...

Why should this work at all?

- What about the 3D geometry of the scene?
- Why aren't we using it?

Image reprojection



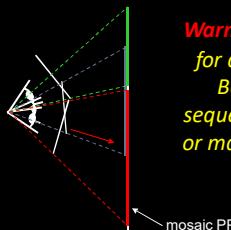
The mosaic has a natural interpretation in 3D:

The images are *reprojected* onto a common plane

The mosaic is formed on this plane.

Source: Steve Seitz

Image reprojection



Warning: This model only holds for angular views up to 180°.
Beyond that need to use sequence that "bends the rays" or map onto a different surface, say, a cylinder.

Mosaics



Obtain a wider angle view by combining multiple images **all of which are taken from the same camera center**.

Image reprojection: Homography

A projective transform is a mapping between any two PPs with the same center of projection

Lines map to lines
So rectangle maps to arbitrary quadrilateral

Called Homography

$$\begin{bmatrix} w \cdot x' \\ w \cdot y' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

p' **H** **p**

Source: Alyosha Efros

Homography

(x_1, y_1) (x_1', y_1')
 (x_2, y_2) (x_2', y_2')
 \vdots \vdots
 (x_n, y_n) (x_n', y_n')

Solving for homographies

$$\mathbf{p}' = \mathbf{H}\mathbf{p} \quad \begin{bmatrix} w \cdot x' \\ w \cdot y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Solving for homographies – non-homogeneous

$$\mathbf{p}' = \mathbf{H}\mathbf{p} \quad \begin{bmatrix} w \cdot x' \\ w \cdot y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Since 8 unknowns, can set scale factor $i=1$.
Set up a system of linear equations $\mathbf{A}\mathbf{h} = \mathbf{b}$ where vector of unknowns

$$\mathbf{h} = [a, b, c, d, e, f, g, h, i]^T$$

Need at least 4 points for 8 eqs, but the more the better...
Solve for \mathbf{h} by $\min \|\mathbf{A}\mathbf{h} - \mathbf{b}\|^2$ using least-squares

Solving for homographies – homogeneous

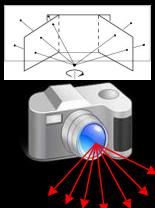
$$\mathbf{p}' = \mathbf{H}\mathbf{p} \quad \begin{bmatrix} w x' \\ w y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Just like we did for the extrinsics, multiply through, and divide out by w . Gives two homogeneous equations per point.
Solve using SVD just like before. This is the cool way.

Apply the Homography

$$p' = \mathbf{H} p \quad \left(\frac{w x'}{w}, \frac{w y'}{w} \right) = (x', y')$$

Mosaics



$$p_i = \begin{bmatrix} -x_i & -y_i & -1 & 0 & 0 & 0 & 0 & x_i x'_i & y_i x'_i & x'_i \\ 0 & 0 & 0 & -x_i & -y_i & -1 & x_i y'_i & y_i y'_i & y'_i \end{bmatrix}$$

$$PH = 0$$

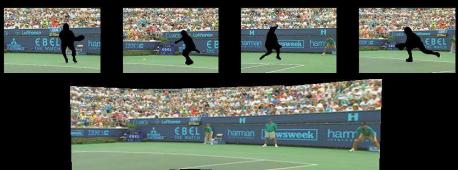
Such as:

$$PH = \begin{bmatrix} -x_1 & -y_1 & -1 & 0 & 0 & 0 & x_1 x'_1 & y_1 x'_1 & x'_1 \\ 0 & 0 & 0 & -x_1 & -y_1 & -1 & x_1 y'_1 & y_1 y'_1 & y'_1 \\ -x_2 & -y_2 & -1 & 0 & 0 & 0 & x_2 x'_2 & y_2 x'_2 & x'_2 \\ 0 & 0 & 0 & -x_2 & -y_2 & -1 & x_2 y'_2 & y_2 y'_2 & y'_2 \\ -x_3 & -y_3 & -1 & 0 & 0 & 0 & x_3 x'_3 & y_3 x'_3 & x'_3 \\ 0 & 0 & 0 & -x_3 & -y_3 & -1 & x_3 y'_3 & y_3 y'_3 & y'_3 \\ -x_4 & -y_4 & -1 & 0 & 0 & 0 & x_4 x'_4 & y_4 x'_4 & x'_4 \\ 0 & 0 & 0 & -x_4 & -y_4 & -1 & x_4 y'_4 & y_4 y'_4 & y'_4 \end{bmatrix} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \\ h_5 \\ h_6 \\ h_7 \\ h_8 \\ h_9 \end{bmatrix} = 0$$

SVD $P = USV^\top$ and select the last singular vector of V as the solution to H .

Mosaics for Video Coding

- Convert masked images into a background sprite for “content-based coding”



Quiz – answer

We said that the transformation between two images taken from the same center of projection is a *homography* H. How many pairs of corresponding points do I need to compute H?

- a) 6
- (b) 4**
- c) 2
- d) 8

Quiz

We said that the transformation between two images taken from the same center of projection is a *homography* H. How many pairs of corresponding points do I need to compute H?

- a) 6
- b) 4
- c) 2
- d) 8

Homographies and 3D planes

Remember this:

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \approx \begin{bmatrix} m_{00} & m_{01} & m_{02} & m_{03} \\ m_{10} & m_{11} & m_{12} & m_{13} \\ m_{20} & m_{21} & m_{22} & m_{23} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

Homographies and 3D planes

- Suppose the 3D points are on a plane:

$$aX + bY + cZ + d = 0$$

Homographies and 3D planes

- On the plane $[a \ b \ c \ d]$ can replace Z:

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \approx \begin{bmatrix} m_{00} & m_{01} & m_{02} & m_{03} \\ m_{10} & m_{11} & m_{12} & m_{13} \\ m_{20} & m_{21} & m_{22} & m_{23} \end{bmatrix} \begin{bmatrix} X \\ Y \\ (aX + bY + d) / (-c) \\ 1 \end{bmatrix}$$

Homographies and 3D planes

- So, can put the Z coefficients into the others:

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = P \begin{bmatrix} X \\ Y \\ (aX + bY + d) / (-c) \\ 1 \end{bmatrix}$$

↑ ↑ ↗

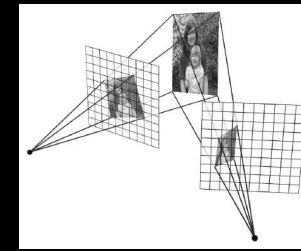
3x3

Homography!

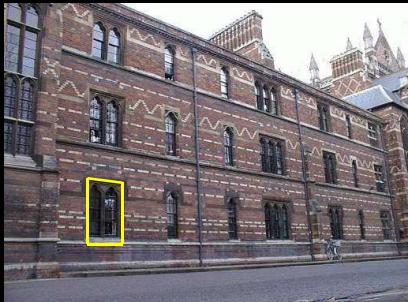
$$H = P \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ -\frac{a}{c} & -\frac{b}{c} & \frac{d}{c} \end{bmatrix}$$

Image reprojection

- Mapping between planes is a homography.
- Whether a plane in the world to the image or between image planes.



Rectifying slanted views



Rectifying slanted views

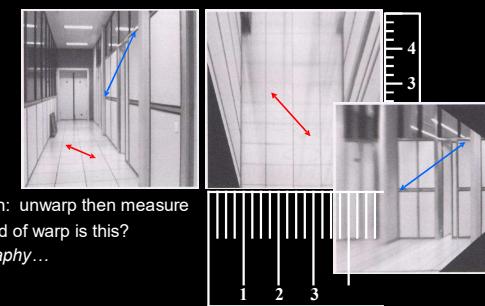


Corrected image (front-to-parallel)

Measuring distances



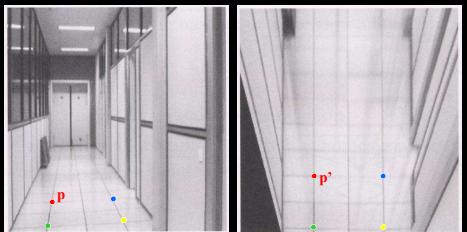
Measurements on planes



Approach: un warp then measure
What kind of warp is this?
Homography...

Image rectification

If there is a planar rectangular grid in the scene you can map it into a rectangular grid in the image...



Some other images of rectangular grids...



Same pixels – via a homography

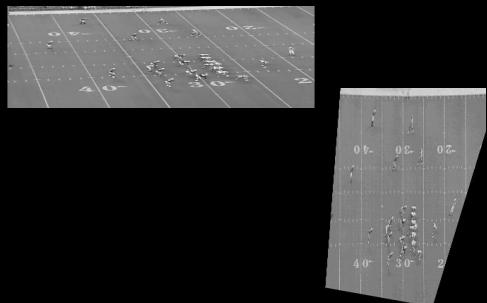
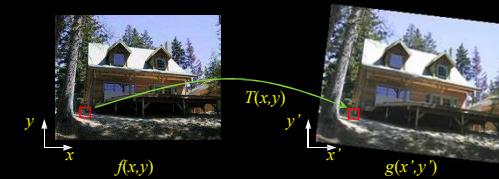


Image warping

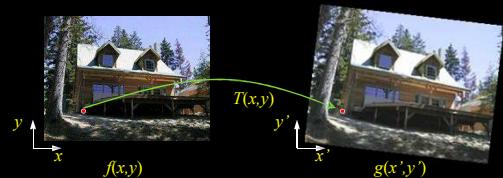
Given a coordinate transform and a source image $f(x,y)$, how do we compute a transformed image $g(x',y') = f(T(x,y))$?



Slide from Alyosha Efros,

Forward warping

Send each pixel $f(x,y)$ to its corresponding location
 $(x',y') = T(x,y)$ in the second image



Q: what if pixel lands “between” two pixels?

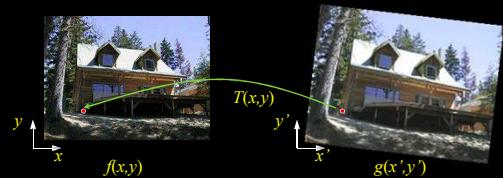
Forward warping

Send each pixel $f(x,y)$ to its corresponding location
 $(x',y') = T(x,y)$ in the second image



Inverse warping

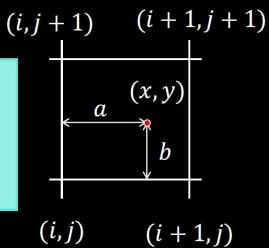
Get each pixel $g(x',y')$ from its corresponding location
 $(x,y) = T^{-1}(x',y')$ in the first image



Q: what if pixel *comes from* “between” two pixels?

Bilinear interpolation

$$\begin{aligned} f(x,y) = & (1-a)(1-b) f[i,j] \\ & +a(1-b) f[i+1,j] \\ & +ab f[i+1,j+1] \\ & +(1-a)b f[i,j+1] \end{aligned}$$



See Matlab (Octave) function `interp2`

Review: How to make a panorama (or mosaic)

Basic Procedure

- Take a sequence of images from the same position
 - Rotate the camera about its optical center
- Compute transformation between second image and first
- Transform the second image to overlap with the first
- Blend the two together to create a mosaic
- (If there are more images, repeat)