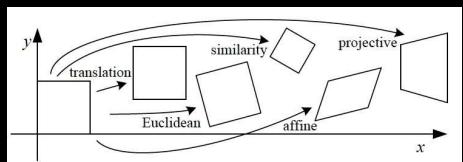


Computer Vision

Introduction to “features”

The basic image point matching problem

- Suppose I have two images related by some transformation. Or have two images of the same object in different positions.
- How to find the transformation of image 1 that would align it with image 2?



Text resources

- Forsyth and Ponce: 5.3-5.4
 - Szeliski also covers this well – Section 4 – 4.1.1

We want *Local⁽¹⁾ Features⁽²⁾*

- Goal: Find points in an image that can be:
 - Found in other images
 - Found precisely – well localized
 - Found reliably – well matched

We want *Local⁽¹⁾ Features⁽²⁾*

Why?

- Want to compute a fundamental matrix to recover geometry
- Robotics/Vision: See how a bunch of points move from one frame to another. Allows computation of how camera moved -> depth -> moving objects
- Build a panorama...

Suppose you want to build a panorama



M. Brown and D. G. Lowe, Recognising Panoramas. ICCV 2003

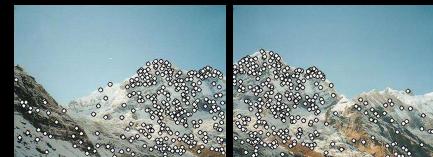
How do we build panorama?

- We need to match (align) images



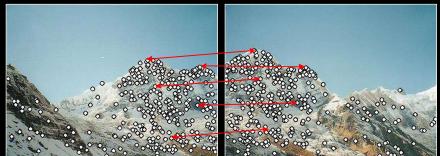
Matching with Features

- Detect features (feature points) in both images



Matching with Features

- Detect features (feature points) in both images
- Match features - find corresponding pairs



Matching with Features

- Detect features (feature points) in both images
- Match features - find corresponding pairs
- Use these pairs to align images



Matching with Features

• Problem 1:

- Detect the same point independently in both images



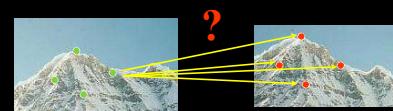
no chance to match!

We need a repeatable detector

Matching with Features

• Problem 2:

- For each point correctly recognize the corresponding one

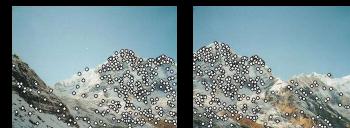


We need a reliable and distinctive **descriptor**

More motivation...

- Feature points are used also for:
 - Image alignment (e.g. homography or fundamental matrix)
 - 3D reconstruction
 - Motion tracking
 - Object recognition
 - Indexing and database retrieval
 - Robot navigation
 - ... other

Characteristics of good features



Characteristics of good features



Repeatability/Precision

- The same feature can be found in several images despite geometric and photometric transformations

Characteristics of good features



Saliency/Matchability

- Each feature has a distinctive description

Characteristics of good features



Compactness and efficiency

- Many fewer features than image pixels

Characteristics of good features

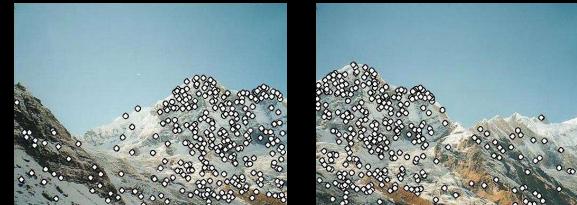


Locality

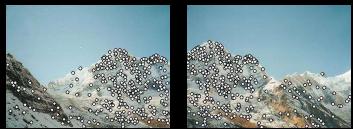
- A feature occupies a relatively small area of the image; robust to clutter and occlusion

Finding corners

Feature points



Characteristics of good features



Repeatability/Precision

- The same feature can be found in several images despite geometric and photometric transformations

Characteristics of good features



Saliency/Matchability

- Each feature has a distinctive description

Characteristics of good features



Compactness and efficiency

- Many fewer features than image pixels

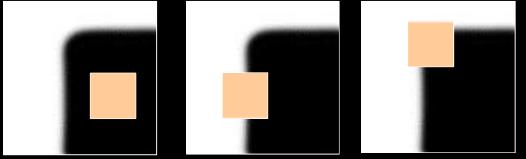
Characteristics of good features



Locality

- A feature occupies a relatively small area of the image; robust to clutter and occlusion

Corner Detection: Basic Idea



"flat" region:
 no change in
 all directions

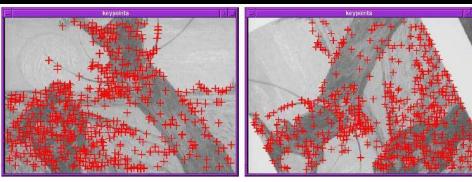
"edge":
 no change
 along the edge
 direction

"corner":
 significant change
 in all directions
 with small shift

Source: A. Efros

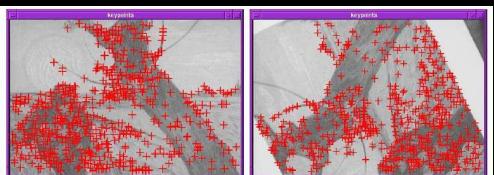
Finding Corners

- Key property: in the region around a corner, image gradient has two or more dominant directions



Finding Corners

C. Harris and M. Stephens. "A Combined Corner and Edge Detector," *Proceedings of the 4th Alvey Vision Conference*: 1988



Corner Detection: Mathematics

Change in appearance for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window function
 Shifted intensity
 Intensity

Window function $w(x,y) =$

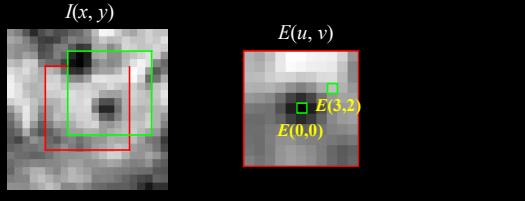
1 in window
0 outside or Gaussian

Source: R. Szeliski

Corner Detection: Mathematics

Change in appearance for the shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u,y+v) - I(x,y)]^2$$



Corner Detection: Mathematics

Change in appearance for the shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u,y+v) - I(x,y)]^2$$

We want to find out how this function behaves for **small** shifts (u,v near 0,0)

Corner Detection: Mathematics

Change in appearance for the shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u,y+v) - I(x,y)]^2$$

Second-order Taylor expansion of $E(u,v)$ about (0,0) (local quadratic approximation for small u,v):

Corner Detection: Mathematics

Change in appearance for the shift $[u,v]$:

$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u,y+v) - I(x,y)]^2$$

$$F(\delta x) \approx F(0) + \delta x \cdot \frac{dF(0)}{dx} + \frac{1}{2} \delta x^2 \cdot \frac{d^2 F(0)}{dx^2}$$

Corner Detection: Mathematics

Change in appearance for the shift $[u, v]$:

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

$$F(\delta x) \approx F(0) + \delta x \cdot \frac{dF(0)}{dx} + \frac{1}{2} \delta x^2 \cdot \frac{d^2 F(0)}{dx^2}$$

$$E(u, v) \approx E(0, 0) + [u \quad v] \begin{bmatrix} E_u(0, 0) \\ E_v(0, 0) \end{bmatrix} + \frac{1}{2} [u \quad v] \begin{bmatrix} E_{uu}(0, 0) & E_{uv}(0, 0) \\ E_{vu}(0, 0) & E_{vv}(0, 0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Second-order Taylor expansion of $E(u, v)$ about (0,0):

$$E(u, v) \approx E(0, 0) + [u \quad v] \begin{bmatrix} E_u(0, 0) \\ E_v(0, 0) \end{bmatrix} + \frac{1}{2} [u \quad v] \begin{bmatrix} E_{uu}(0, 0) & E_{uv}(0, 0) \\ E_{vu}(0, 0) & E_{vv}(0, 0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Second-order Taylor expansion of $E(u, v)$ about (0,0):

$$E(u, v) \approx E(0, 0) + [u \quad v] \begin{bmatrix} E_u(0, 0) \\ E_v(0, 0) \end{bmatrix} + \frac{1}{2} [u \quad v] \begin{bmatrix} E_{uu}(0, 0) & E_{uv}(0, 0) \\ E_{vu}(0, 0) & E_{vv}(0, 0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

Need these derivatives...

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Second-order Taylor expansion of $E(u, v)$ about (0,0):

$$E_u(u, v) = \sum_{x, y} 2w(x, y) [I(x + u, y + v) - I(x, y)] I_y(x + u, y + v)$$

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Second-order Taylor expansion of $E(u, v)$ about (0,0):

$$E_{uu}(u, v) = \sum_{x,y} 2w(x, y) I_x(x+u, y+v) I_x(x+u, y+v)$$

$$+ \sum_{x,y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_{xx}(x+u, y+v)$$

$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Second-order Taylor expansion of $E(u, v)$ about (0,0):

$$E_{vv}(u, v) = \sum_{x,y} 2w(x, y) I_y(x+u, y+v) I_y(x+u, y+v)$$

$$+ \sum_{x,y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_{yy}(x+u, y+v)$$

Second-order Taylor expansion of $E(u, v)$ about (0,0):

$$E(u, v) \approx E(0, 0) + [u - v] \begin{bmatrix} E_x(0, 0) \\ E_{xx}(0, 0) \end{bmatrix} + \frac{1}{2}[u - v] \begin{bmatrix} E_{xx}(0, 0) & E_{xy}(0, 0) \\ E_{yx}(0, 0) & E_{yy}(0, 0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E_x(u, v) = \sum_{x,y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_x(x+u, y+v)$$

$$E_{uv}(u, v) = \sum_{x,y} 2w(x, y) I_x(x+u, y+v) I_y(x+u, y+v)$$

$$+ \sum_{x,y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_{xy}(x+u, y+v)$$

$$E_{yy}(u, v) = \sum_{x,y} 2w(x, y) I_y(x+u, y+v) I_y(x+u, y+v)$$

$$+ \sum_{x,y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_{yy}(x+u, y+v)$$

Evaluate E and its derivatives at (0,0):

$$E(u, v) \approx E(0, 0) + [u - v] \begin{bmatrix} E_x(0, 0) \\ E_{xx}(0, 0) \end{bmatrix} + \frac{1}{2}[u - v] \begin{bmatrix} E_{xx}(0, 0) & E_{xy}(0, 0) \\ E_{yx}(0, 0) & E_{yy}(0, 0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E_x(0, 0) = \sum_{x,y} 2w(x, y) [I(x, y) - I(x, y)] = 0$$

$$E_{xx}(0, 0) = \sum_{x,y} 2w(x, y) I_x(x, y) I_x(x, y) = 0$$

$$E_{yy}(0, 0) = \sum_{x,y} 2w(x, y) I_y(x, y) I_y(x, y) = 0$$

$$E_{xy}(0, 0) = \sum_{x,y} 2w(x, y) I_x(x, y) I_y(x, y) = 0$$

$$E_{yx}(0, 0) = \sum_{x,y} 2w(x, y) I_y(x, y) I_x(x, y) = 0$$

Second-order Taylor expansion of $E(u,v)$ about (0,0):

$$E(u,v) \approx E(0,0) + [u-v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2}[u-v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{vu}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0,0) = 0 \quad E_{uu}(0,0) = \sum_{x,y} 2w(x,y)I_x(x,y)I_x(x,y)$$

$$E_u(0,0) = 0 \quad E_{vu}(0,0) = \sum_{x,y} 2w(x,y)I_x(x,y)I_y(x,y)$$

$$E_v(0,0) = 0 \quad E_{uv}(0,0) = \sum_{x,y} 2w(x,y)I_y(x,y)I_x(x,y)$$

Second-order Taylor expansion of $E(u,v)$ about (0,0):

$$E(u,v) \approx [u-v] \begin{bmatrix} \sum_{x,y} w(x,y)I_x^+(x,y) & \sum_{x,y} w(x,y)I_x(x,y)I_y(x,y) \\ \sum_{x,y} w(x,y)I_x(x,y)I_y(x,y) & \sum_{x,y} w(x,y)I_y^+(x,y) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0,0) = 0 \quad E_{uu}(0,0) = \sum_{x,y} 2w(x,y)I_x(x,y)I_x(x,y)$$

$$E_u(0,0) = 0 \quad E_{vv}(0,0) = \sum_{x,y} 2w(x,y)I_y(x,y)I_y(x,y)$$

$$E_v(0,0) = 0 \quad E_{uv}(0,0) = \sum_{x,y} 2w(x,y)I_x(x,y)I_y(x,y)$$

Corner Detection: Mathematics

The quadratic approximation simplifies to

$$E(u,v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a *second moment matrix* computed from image derivatives:

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

The second moment matrix M :

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Each product is
a rank 1 2x2

Can be written (without the weight):

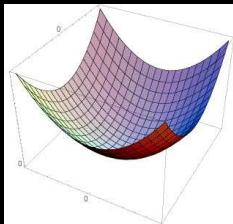
$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \left(\begin{bmatrix} I_x \\ I_y \end{bmatrix} \begin{bmatrix} I_x & I_y \end{bmatrix} \right) = \sum \nabla I (\nabla I)^T$$

Interpreting the second moment matrix

The surface $E(u,v)$ is locally approximated by a quadratic form.

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

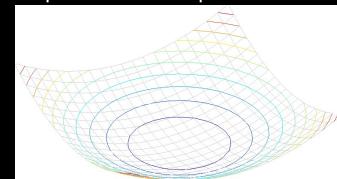


Interpreting the second moment matrix

Consider a constant “slice” of $E(u, v)$:

$$\Sigma I_x^2 u^2 + 2 \Sigma I_x I_y u v + \Sigma I_y^2 v^2 = k$$

This is the equation of an ellipse.



Interpreting the second moment matrix

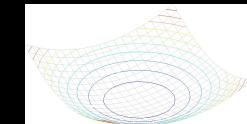
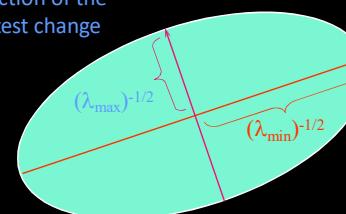
First, consider the axis-aligned case where gradients are either horizontal or vertical

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

If either λ is close to 0, then this is **not** a corner, so look for locations where both are large.

Interpreting the second moment matrix

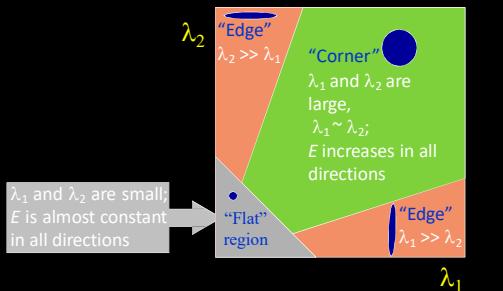
direction of the fastest change



direction of the slowest change

Interpreting the eigenvalues

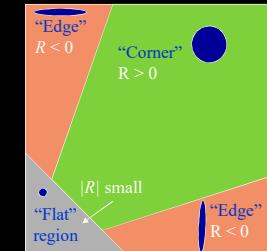
Classification of image points using eigenvalues of M :



Harris corner response function

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

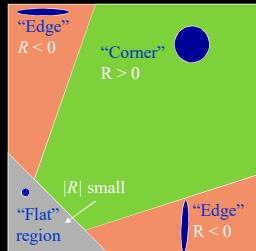
α : constant (0.04 to 0.06)



Harris corner response function

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

R depends only on eigenvalues of M , but don't compute them (no sqrt, so really fast even in the '80s).



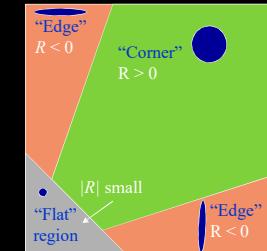
Harris corner response function

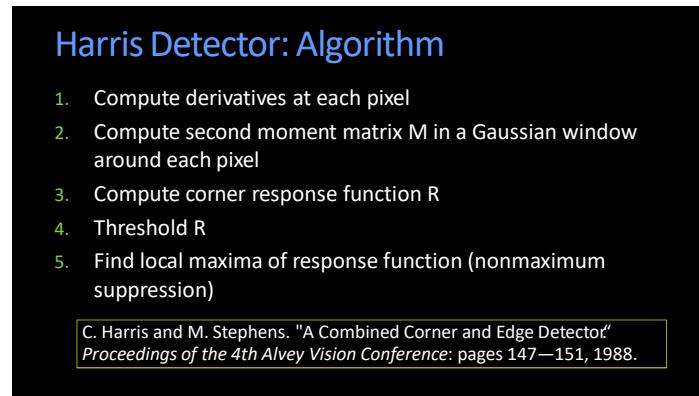
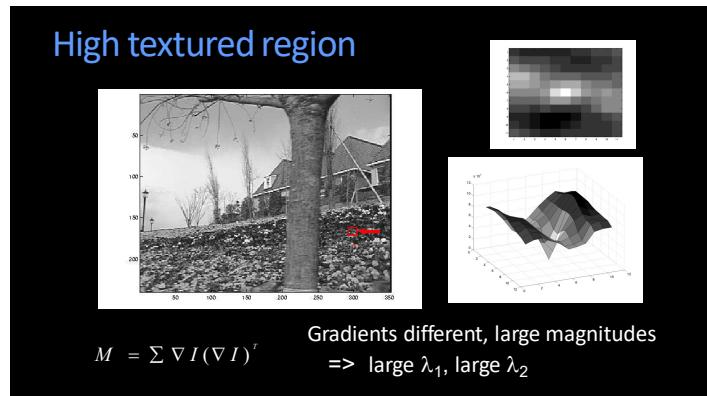
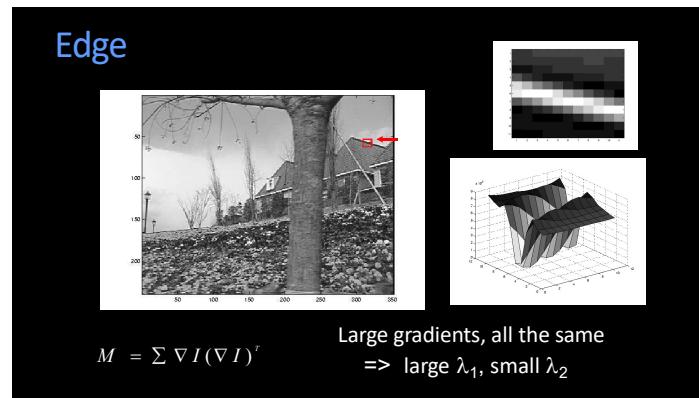
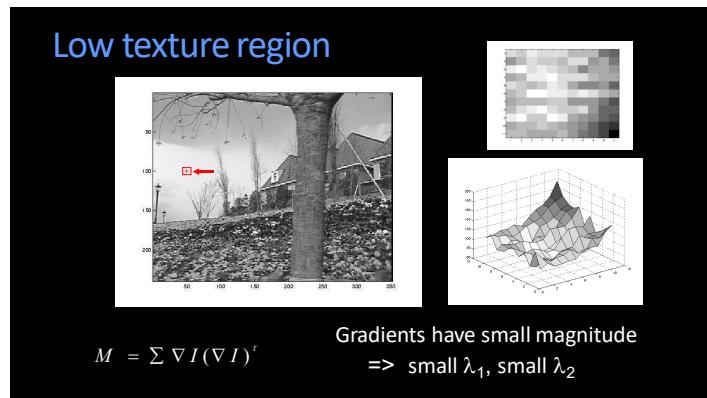
$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

R is large for a corner

R is negative with large magnitude for an edge

$|R|$ is small for a flat region





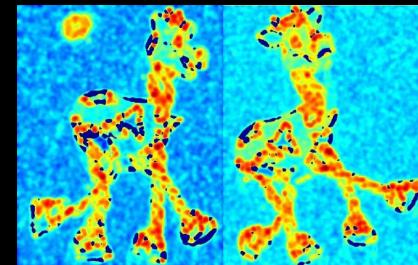
21.10.2025

Harris Detector: Workflow



Harris Detector: Workflow

Compute corner response R



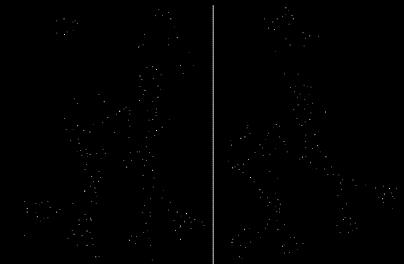
Harris Detector: Workflow

Find points with large corner response: $R > \text{threshold}$



Harris Detector: Workflow

Take only the points of local maxima of R



Harris Detector: Workflow

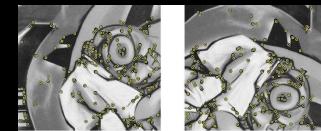


Other corners:

Shi-Tomasi '94:

"Cornerness" = $\min(\lambda_1, \lambda_2)$ Find local maximums
`cvGoodFeaturesToTrack(...)`

Reportedly better for regions undergoing affine deformations



Other corners:

- Brown, M., Szeliski, R., and Winder, S. (2005):

$$\frac{\det M}{\text{tr } M} = \frac{\lambda_0 \lambda_1}{\lambda_0 + \lambda_1}$$

- There are others...

Computer Vision

Scale invariance

Recall: Corner Detection - Basic Idea

"flat" region:
no change in
all directions

"edge":
no change
along the edge
direction

"corner":
significant change
in all directions
with small shift

Source: A. Efros

Corner Detection: Mathematics

Change in appearance for the shift $[u, v]$:

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

$I(x, y)$

$E(u, v)$

$E(3,2)$

$E(0,0)$

Corner Detection: Mathematics

The quadratic approximation simplifies to

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a *second moment matrix* computed from image derivatives:

$$M = \sum_{x, y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

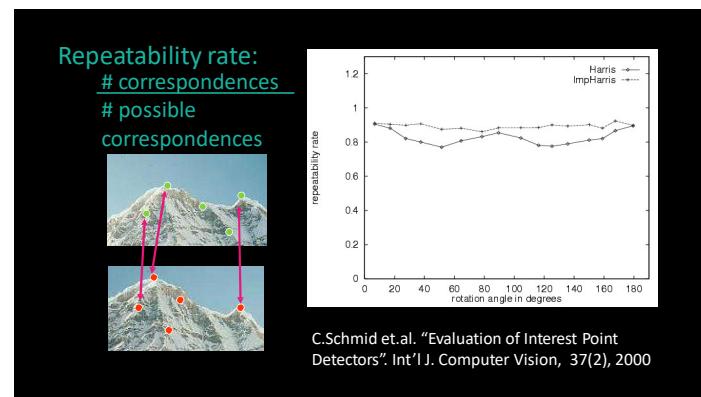
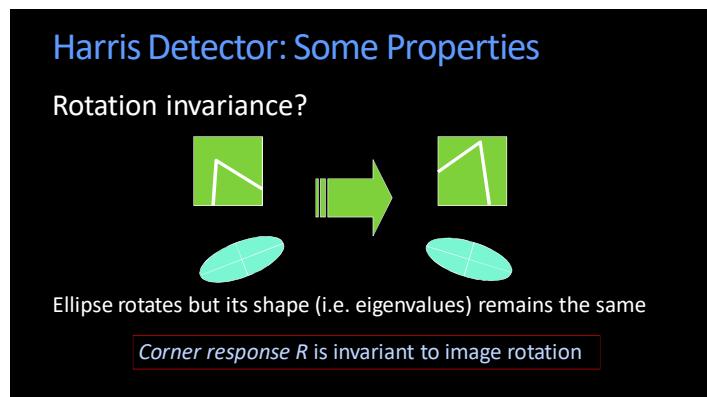
Harris corner response function

$$R = \det(M) - \alpha \text{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

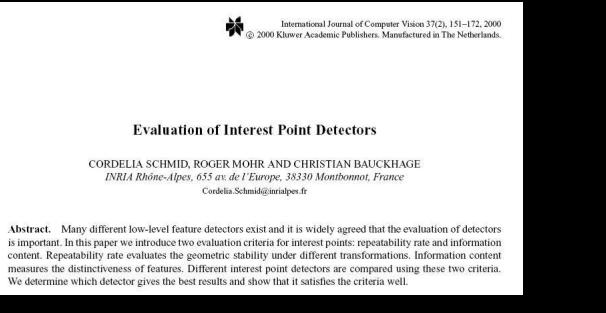
R is large for a corner

R is negative with large magnitude for an edge

$|R|$ is small for a flat region



Evaluation plots are from this paper



Evaluation of Interest Point Detectors

CORDELIA SCHMID, ROGER MOHR AND CHRISTIAN BAUCKHAGE
INRIA Rhône-Alpes, 655 av de l'Europe, 38330 Montbonnot, France
Cordeelia.Schmid@inrialpes.fr

Abstract. Many different low-level feature detectors exist and it is widely agreed that the evaluation of detectors is important. In this paper we introduce two evaluation criteria for interest points: repeatability rate and information content. Repeatability rate evaluates the geometric stability under different transformations. Information content measures the distinctiveness of features. Different interest point detectors are compared using these two criteria. We determine which detector gives the best results and show that it satisfies the criteria well.

Harris Detector: Some Properties

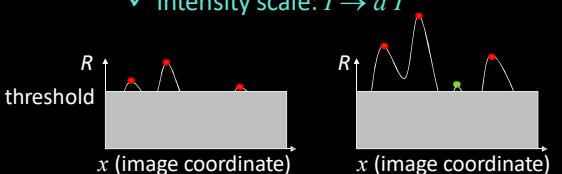
- Invariance to image intensity change?

Harris Detector: Some Properties

- Mostly invariant to additive and multiplicative intensity changes (threshold issue for multiplicative)
 - ✓ Only derivatives are used => invariance to intensity shift $I \rightarrow I + b$

Harris Detector: Some Properties

- Mostly invariant to additive and multiplicative intensity changes (threshold issue for multiplicative)
 - ✓ Intensity scale: $I \rightarrow a I$

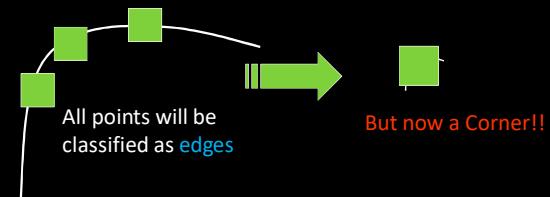


Harris Detector: Some Properties

- Invariant to image scale?

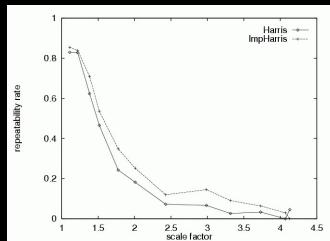
Harris Detector: Some Properties

- Not invariant to *image scale!*



Harris Detector: Some Properties

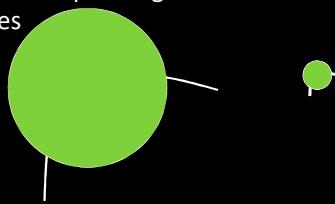
Not invariant to image scale:



IF we want scale invariance...

Scale Invariant Detection

- Consider regions (e.g. circles) of different sizes around a point
- Regions of corresponding sizes will look the same in both images



Scale Invariant Detection

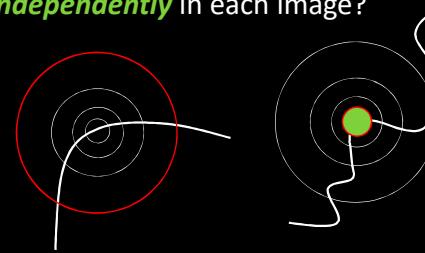
Solution:

- Design a function on the region (circle), which is “scale invariant” - not affected by the size but will be the same for “corresponding regions,” even if they are at different sizes/scales.

Example: Average intensity. For corresponding regions (even of different sizes) it will be the same.

Scale Invariant Detection

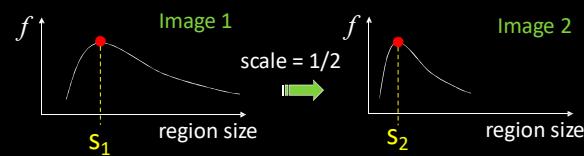
- The problem: how do we choose corresponding circles **independently** in each image?

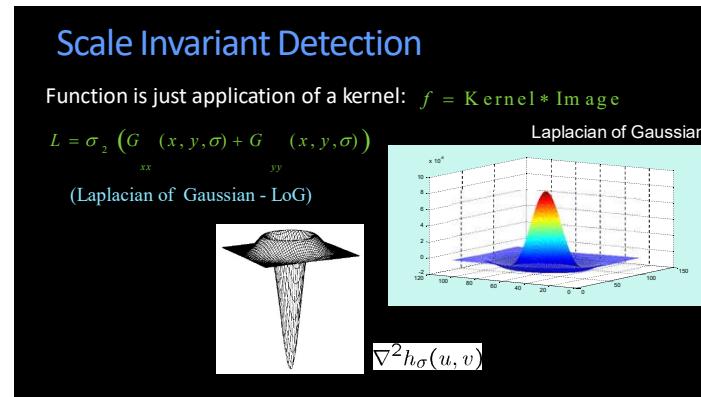
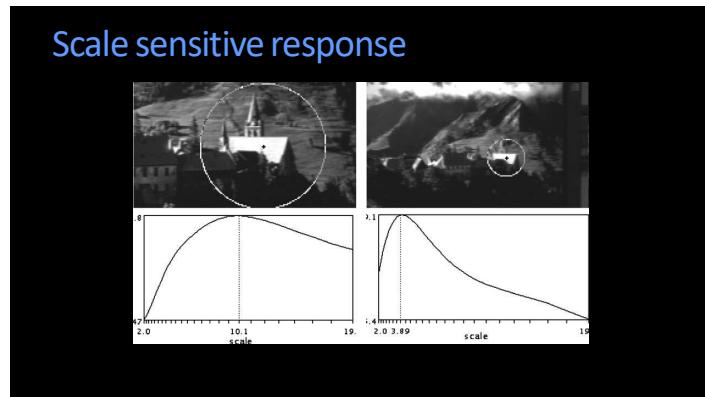
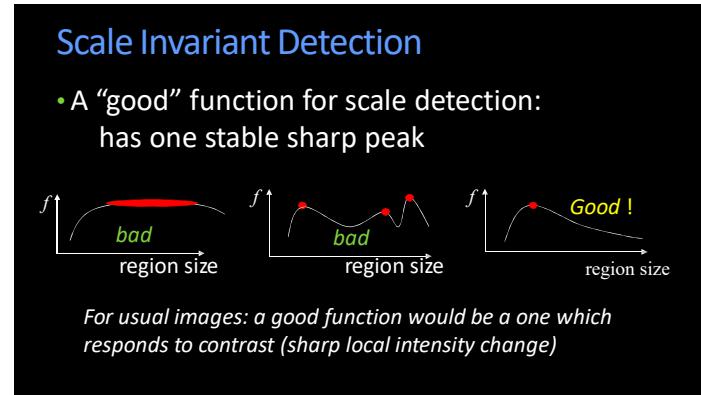
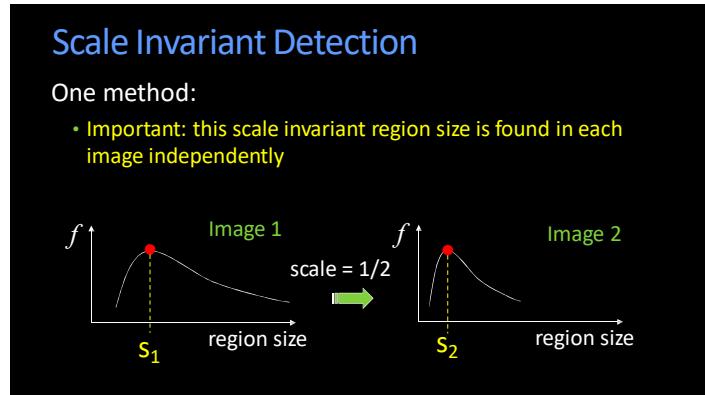


Scale Invariant Detection

One method:

- At a point, compute the scale invariant function over different size neighborhoods (different scales).
- Choose the scale for each image at which the function is a maximum





Scale Invariant Detection

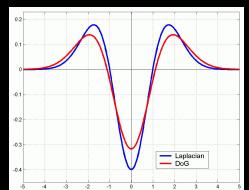
Function is just application of a kernel: $f = \text{Kernel} * \text{Image}$

$$L = \sigma_2 \left(G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

(Laplacian of Gaussian - LoG)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

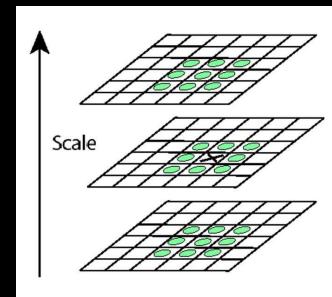
(Difference of Gaussians)



Note: both kernels are invariant to scale and rotation

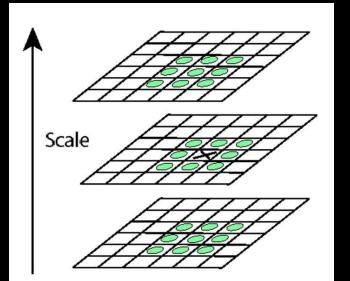
Key point localization

- General idea: find robust extremum (maximum or minimum) both in space and in scale.



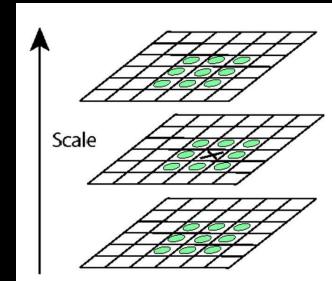
Key point localization

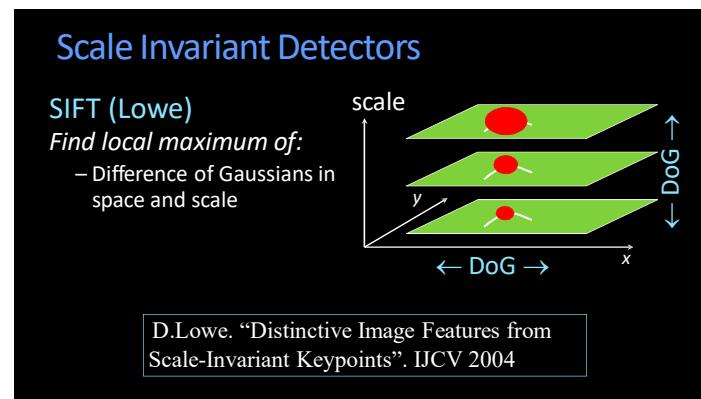
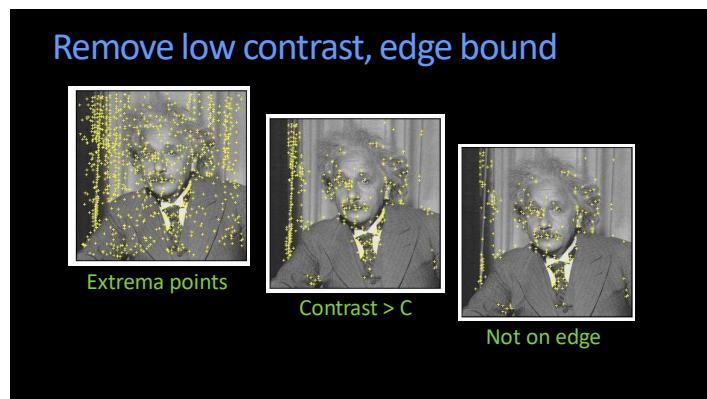
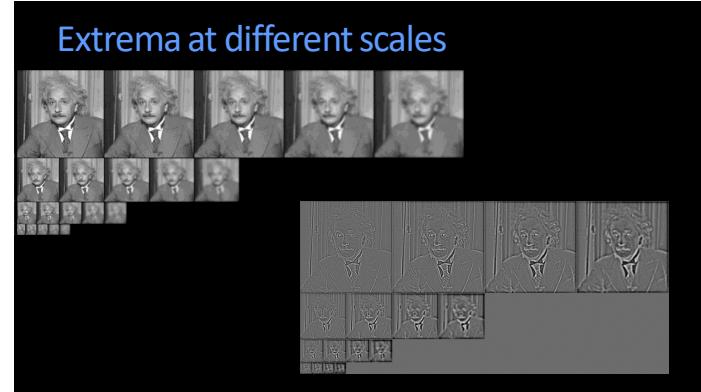
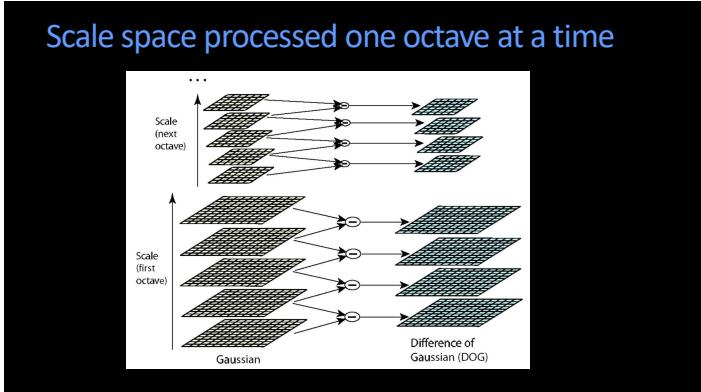
- SIFT: Scale Invariant Feature Transform
- Specific suggestion: use DoG pyramid to find maximum values (remember edge detection?) – then eliminate “edges” and pick only corners.



Key point localization

(Each point is compared to its 8 neighbors in the current image and 9 neighbors each in the scales above and below.)



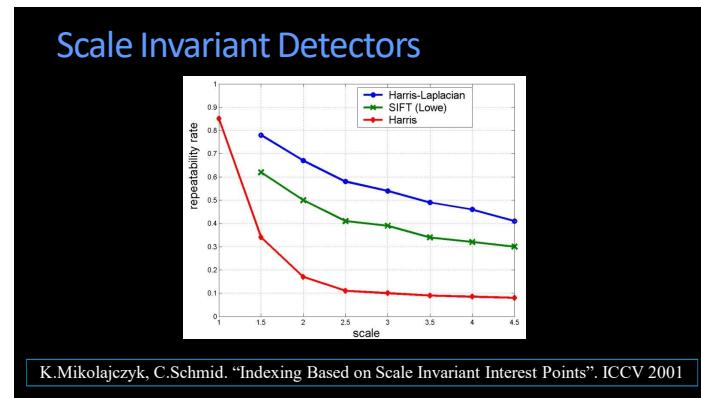


Scale Invariant Detectors

Harris-Laplacian
Find local maximum of:

- Harris corner detector in space (image coordinates)
- Laplacian in scale

K.Mikolajczyk, C.Schmid. "Indexing Based on Scale Invariant Interest Points". ICCV 2001



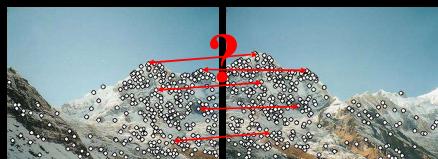
Scale Invariance: Summary

- Given: Two images of the same scene with a **large** scale difference between them
- Goal: Find the same interest points independently in each image
- Solution: Search for maxima of suitable functions in scale and in space (over the image)



Point Descriptors

- Next question: How to match them?



Point Descriptors

- We need to describe them – a “*descriptor*”

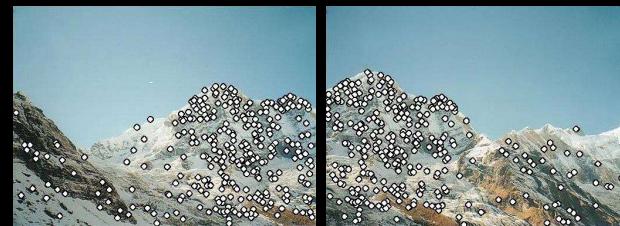


...Next!

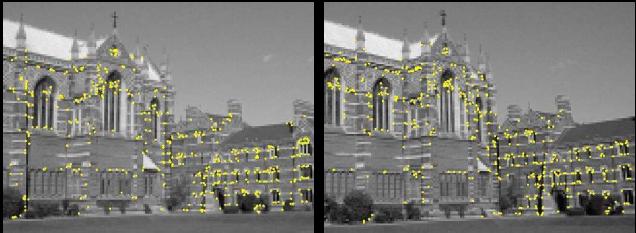
SIFT descriptor

Point Descriptors

- Last time: How to detect interest points



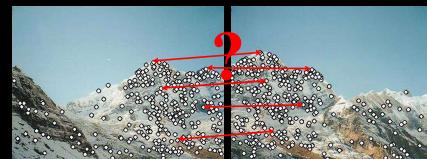
Harris detector



Interest points extracted with Harris (~ 500 points)

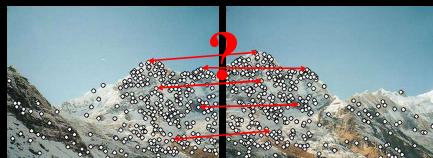
Point Descriptors

- Now: How to match them?



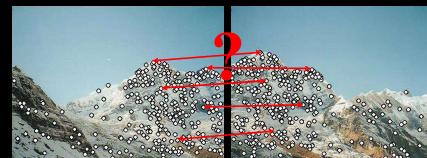
Point Descriptors

- We need to describe them – a “*descriptor*”



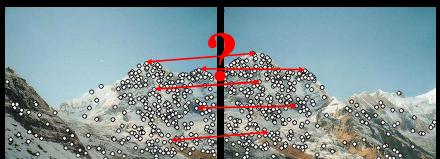
Criteria for Point Descriptors

- We want the descriptors to be the (almost) same in both image – *invariant*.



Criteria for Point Descriptors

- We also need the descriptors to be *distinctive*.



Simple solution?

- Harris gives good detection – and we also know the scale.
- Why not just use correlation to check the match of the window around the feature in image 1 with every feature in image 2?

Simple solution? *Not so good!*

- Not so good because:
 - Correlation is not rotation invariant - why do we want this?
 - Correlation is sensitive to photometric changes.
 - Normalized correlation is sensitive to non-linear photometric changes and even slight geometric ones.
 - Could be slow – check all features against all features

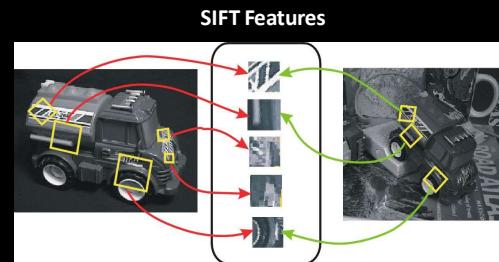
SIFT: Scale Invariant Feature Detection

- Motivation: The Harris operator was not invariant to scale and correlation was not invariant to rotation.
- For better image matching, Lowe's goals were:
 - To develop an interest operator – a *detector* – that is invariant to scale and rotation.
 - Also: create a *descriptor* that was robust to the variations corresponding to typical viewing conditions. *The descriptor is the most-used part of SIFT.*

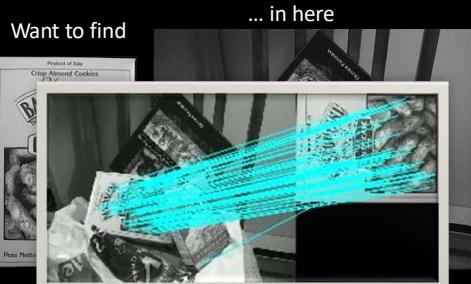
Idea of SIFT

- Image content is represented by a constellation of local features that are invariant to translation, rotation, scale, and other imaging parameters

Idea of SIFT



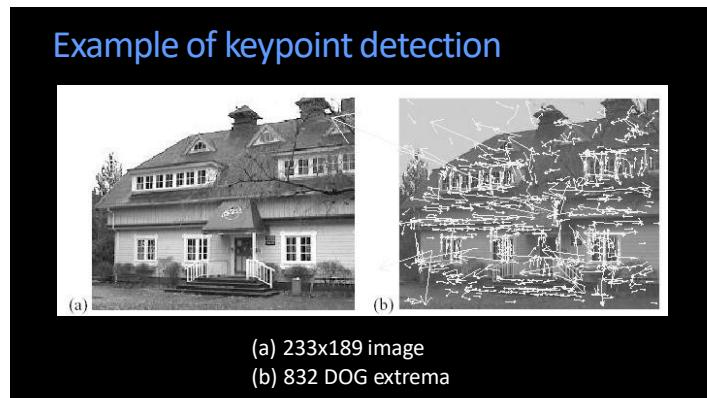
Another version of the problem...



Overall SIFT Procedure

- Scale-space extrema detection
- Keypoint localization
- Orientation assignment
- Keypoint description

Or use Harris-Laplace or other method

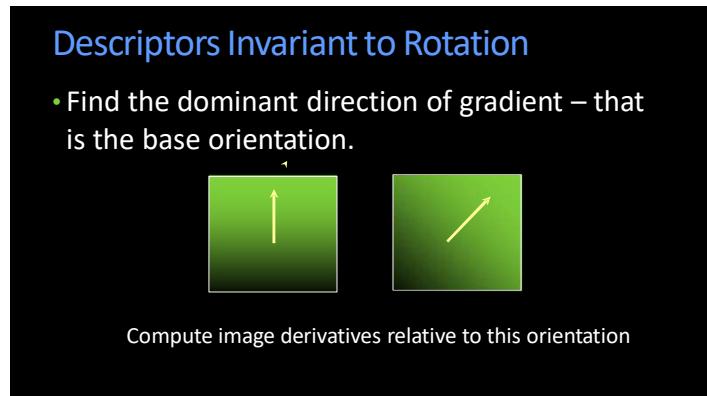


Overall SIFT Procedure

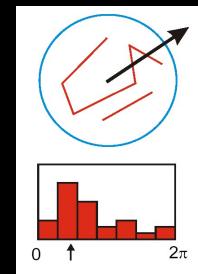
1. Scale-space extrema detection
2. Keypoint localization
3. Orientation assignment

Compute best orientation(s) for each keypoint region.
4. Keypoint description

Use local image gradients at selected scale and rotation to describe each keypoint region.

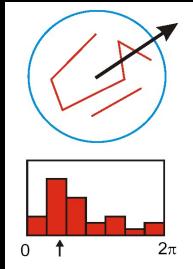


Orientation assignment



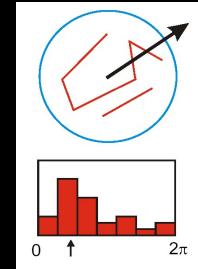
- Create histogram of local gradient directions at *selected* scale – 36 bins

Orientation assignment



- Assign canonical orientation at peak of smoothed histogram

Orientation assignment



- Each **keypoint** now specifies stable 2D coordinates ($x, y, scale, orientation$) – invariant to those.

4. Keypoint Descriptors

- Next is to compute a descriptor for the local image region about each keypoint that is:
 - Highly **distinctive**
 - As **invariant** as possible to variations such as changes in viewpoint and illumination

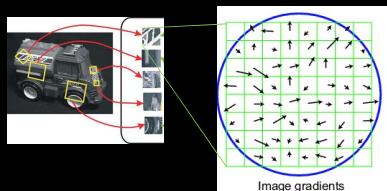
But first... normalization...

- Rotate the window to standard orientation
- Scale the window size based on the scale at which the point was found.

SIFT vector formation

Compute a feature vector based upon:

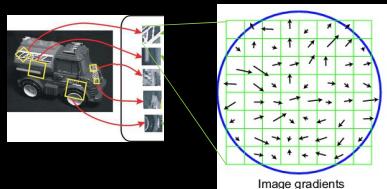
- histograms of gradients



SIFT vector formation

Compute a feature vector based upon:

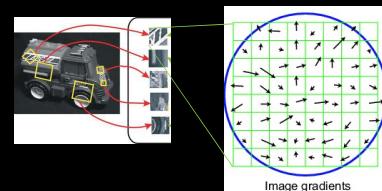
- weighted by the magnitude of the gradient



SIFT vector formation

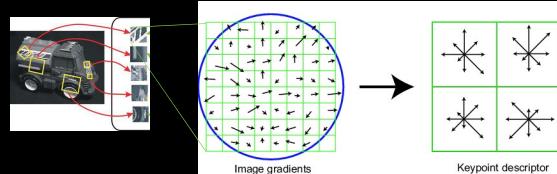
Compute a feature vector based upon:

- weighted by a centered Gaussian,

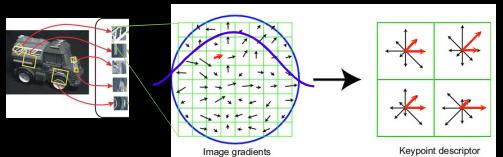


SIFT vector formation

showing only
2x2 here but
it really is 4x4

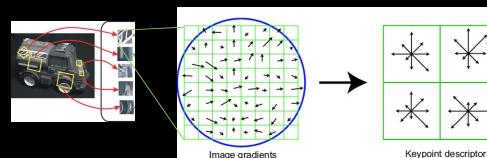


Ensure smoothness



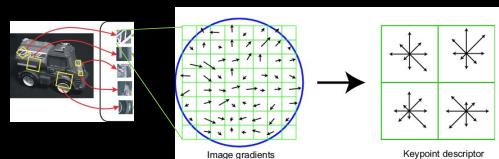
Reduce effect of illumination

- Clip gradient magnitudes to avoid excessive influence of high gradients
 - after rotation normalization, clamp gradients >0.2



Reduce effect of illumination

- 128-dim vector normalized to magnitude 1.0



Evaluating the SIFT descriptors

- Database images were subjected to rotation, scaling, affine stretch, brightness and contrast changes, and added noise.
- Feature point detectors and descriptors were compared before and after the distortions.
- Mostly looking for stability with respect to change.

Sensitivity to parameters

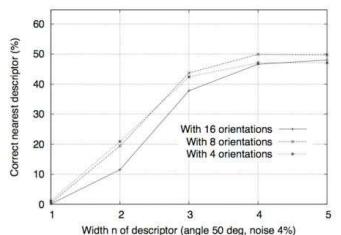
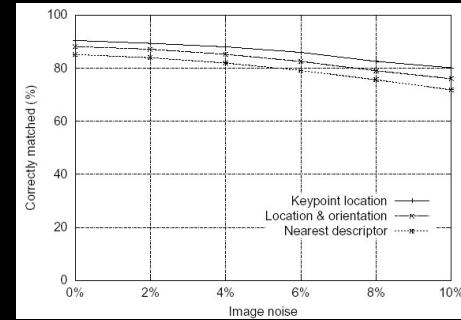


Figure 8: This graph shows the percent of keypoints giving the correct match to a database of 40,000 keypoints as a function of width of the $n \times n$ keypoint descriptor and the number of orientations in each histogram. The graph is computed for images with affine viewpoint change of 50 degrees and addition of 4% noise.

Feature stability to noise

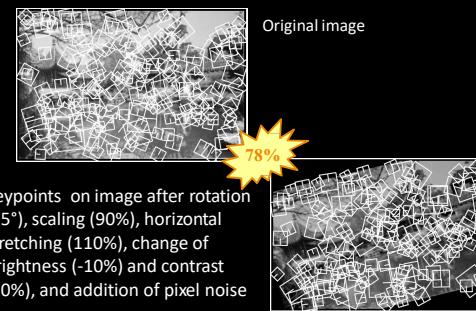


Experimental results - summary

Image transformation	Location and scale match	Orientation match
Decrease contrast by 1.2	89.0 %	86.6 %
Decrease intensity by 0.2	88.5 %	85.9 %
Rotate by 20°	85.4 %	81.0 %
Scale by 0.7	85.1 %	80.3 %
Stretch by 1.2	83.5 %	76.1 %
Stretch by 1.5	77.7 %	65.0 %
Add 10% pixel noise	90.3 %	88.4 %
All previous	78.6 %	71.8 %

20 different images, around 15,000 keys

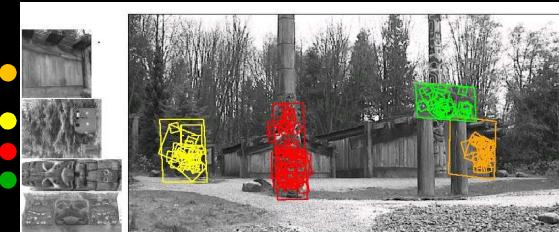
Experimental results



SIFT matching object pieces (for location)



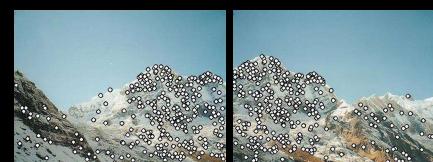
SIFT matching object pieces (for location)



Matching feature points (a little)

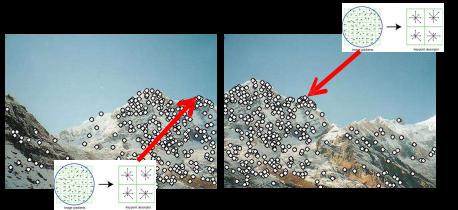
Feature Points

- We know how to detect points



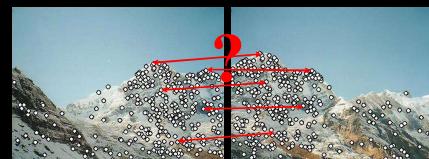
Feature Points

- We know how to describe them



Feature Points

- Next question: How to match them?



How to match feature points?

- Could just do nearest-neighbor search
 - [OMS students: You will!]
- But that's really expensive...SIFT tests have 10,000's of points!

Nearest-neighbor matching to feature database

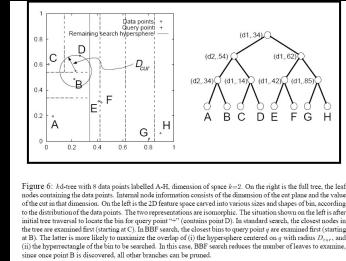
- Better: Hypotheses are generated by *approximate nearest neighbor* matching of each feature to vectors in the database
 - SIFT uses best-bin-first (Beis & Lowe, 97) modification to k-d tree algorithm
 - Use heap data structure to identify bins in order by their distance from query point

Nearest-neighbor matching to feature database

- Result: Can give speedup by factor of **100-1000** while finding nearest neighbor (of interest) 95% of the time

Nearest neighbor techniques

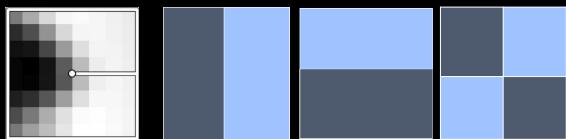
- k -D tree and
- Best Bin First (BBF)



Indexing Without Invariants in 3D Object Recognition, Beis and Lowe, PAMI'99

Wavelet-based hashing

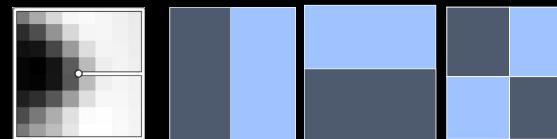
Compute a short (3-vector) descriptor from the neighborhood using a Haar “wavelet”



[Brown, Szeliski, Winder, CVPR'2005]

Wavelet-based hashing

Quantize each value into 10 (overlapping) bins (10^3 total entries)



[Brown, Szeliski, Winder, CVPR'2005]

3D Object Recognition

Train:

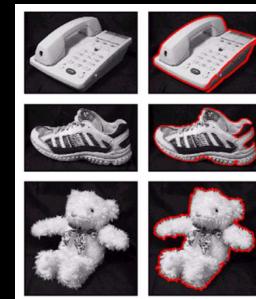
1. Extract outlines with background subtraction
2. Compute “keypoints” – interest points and descriptors.



3D Object Recognition

Test:

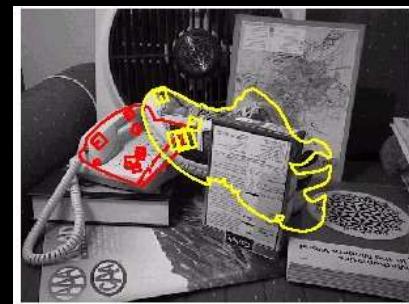
1. Find possible matches.
2. Search for consistent solution – such as *affine*.
(*How many points?!?!*)



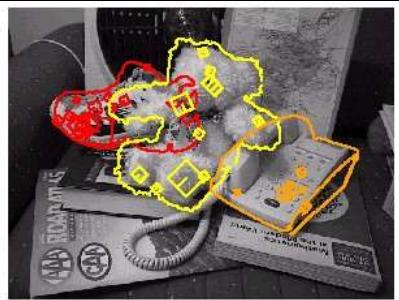
Results



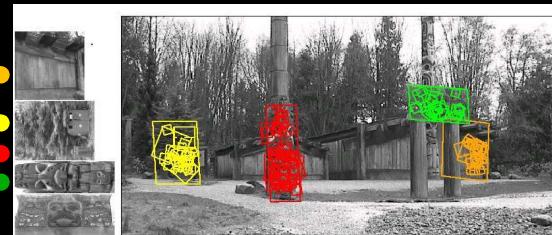
Recognition under occlusion



Recognition under occlusion



Locating object pieces



(From last lesson)

SIFT in Sony Aibo (Evolution Robotics)

SIFT usage:

- Recognize charging station
- Communicate with visual cards

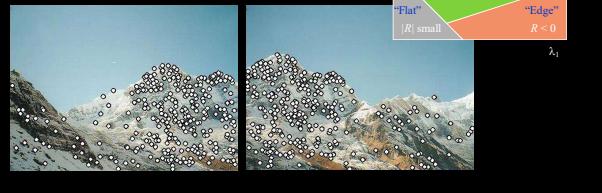


Robust error functions

Feature-based alignment to find transforms

Overall strategy:

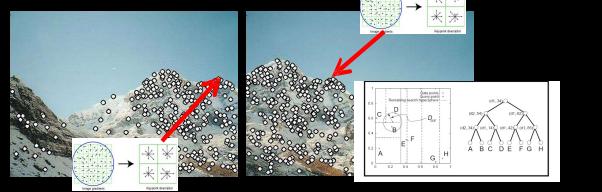
1. Compute features – detect and describe



Feature-based alignment to find transforms

Overall strategy:

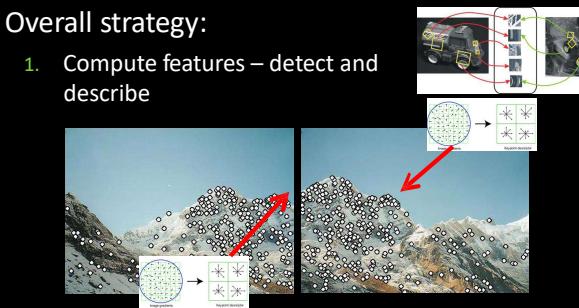
2. Find some useful matches:
Kd-tree, Best-Bin, Hashing



Feature-based alignment to find transforms

Overall strategy:

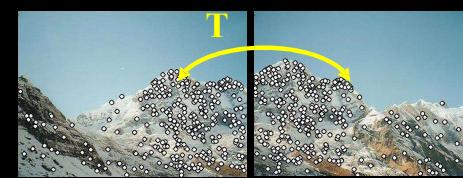
1. Compute features – detect and describe



Feature-based alignment to find transforms

Overall strategy:

3. Compute and apply the best transformation:
e.g. affine, translation, or homography



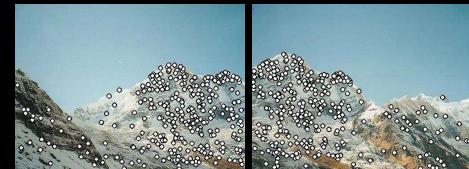
Feature-based alignment to find transforms

Overall strategy:

3. Compute and **apply** the best transformation:
e.g. *affine, translation, or homography*



Feature-based alignment algorithm



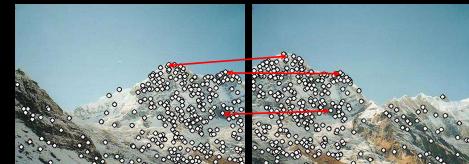
1. Extract features

Feature-based alignment algorithm



2. Compute ***putative* matches** – e.g. “closest descriptor”
Kd-tree, best bin, etc...

Feature-based alignment algorithm



3. Loop until happy:
 - *Hypothesize* transformation T from some matches
 - *Verify* transformation (search for other matches consistent with T) – mark best

Feature-based alignment algorithm



4. Apply best transformation.

How to get “putative” matches?

Feature matching

- Exhaustive search
- Hashing
- Nearest neighbor techniques

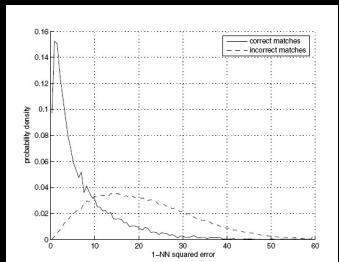
.... but these give the best match. How do we know it's a good one?

Feature-space outlier rejection

- Let's not match all features, but only these that have “similar enough” matches?
- How can we do it?
 - $SSD(patch1,patch2) < threshold$
 - How to set threshold?

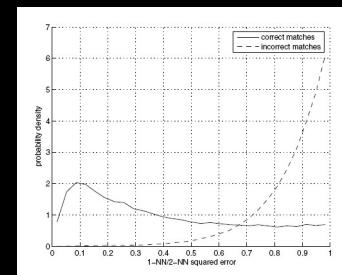
Feature-space outlier rejection

- How to set threshold?



A better way [Lowe, 1999]:

- 1-NN: SSD of the closest match
- 2-NN: SSD of the second-closest match
- Look at how much better 1-NN is than 2-NN, e.g. 1-NN/2-NN

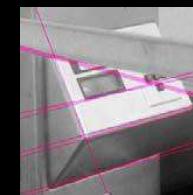


Feature matching

- Exhaustive search
- Hashing
- Nearest neighbor techniques
- But...remember the distinctive vs invariant competition? Implies:
- Problem: Even when pick best match, still lots (and lots) of wrong matches – “outliers”. What should we do?

Model Fitting

- Choose a parametric model to represent a set of features – *remember this???*



simple model: lines



simple model: circles

Fitting: Issues

Case study: Line detection

- **Noise** in the measured feature locations
- **Extraneous data**: clutter (outliers), multiple lines
- **Missing data**: occlusions



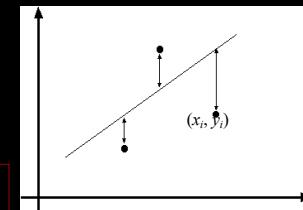
Slide: S. Lazebnik

RANSAC

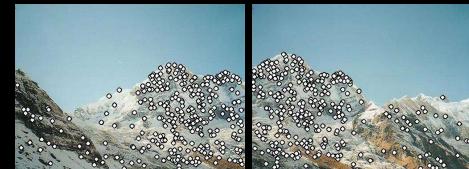
Typical least squares line fitting

- Data: $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation: $y_i = m x_i + b$
- Find (m, b) to minimize:

$$E = \sum_{i=1}^n (y_i - m x_i - b)^2$$



Feature-based alignment algorithm



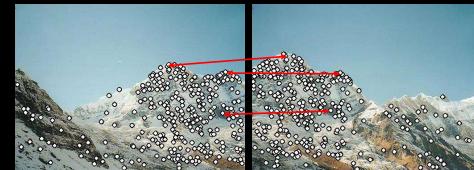
1. Extract features

Feature-based alignment algorithm



2. Compute ***putative*** matches – e.g. “closest descriptor”
Kd-tree, best bin, etc...

Feature-based alignment algorithm

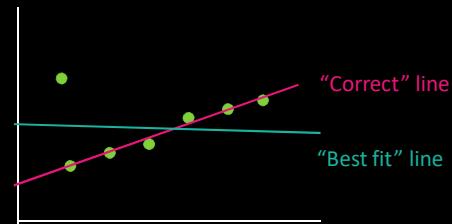


3. Loop until happy:
 - Hypothesize transformation T from some matches
 - Verify transformation (search for other matches ***consistent*** with T) – mark best

“Find consistent matches”?

- Some “best” matches are correct
- Some are not. And the “not” are not part of any other consistent match...
- Need to find the right ones so can compute the pose/transform/fundamental... ***the model***.
- Today: Random Sample Consensus (RANSAC)

Simple Example: Fitting a line



RANSAC: Main idea

- Fitting a line (model) is easy if we know which points belong and which do not. (duh...)
- If we had a proposed line (model), we could probably guess which points belong to that line (model): *inliers*.
- **R**AN**D** **S**AMP**E** **C**ONSENSUS: randomly pick some points to define your line (model). Repeat enough times until you find a good line (model) – one with many inliers.
- Fischler & Bolles 1981 – Copes with a large proportion of outliers

RANSAC

Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using sample
3. **Score** by the fraction of *inliers* within a preset threshold of the model
4. **Repeat** 1-3 until the best model is found with high confidence

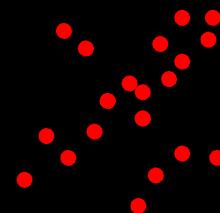
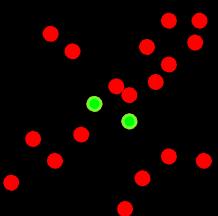


Illustration by Savarese

RANSAC

Algorithm:

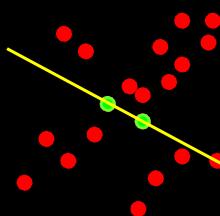
1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using sample
3. **Score** by the fraction of *inliers* within a preset threshold of the model
4. **Repeat** 1-3 until the best model is found with high confidence



RANSAC

Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using sample
3. **Score** by the fraction of *inliers* within a preset threshold of the model
4. **Repeat** 1-3 until the best model is found with high confidence



RANSAC

Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using sample
3. **Score** by the fraction of *inliers* within a preset threshold of the model
4. **Repeat** 1-3 until the best model is found with high confidence

$|C_i| = 6$

RANSAC

Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using sample
3. **Score** by the fraction of *inliers* within a preset threshold of the model
4. **Repeat** 1-3 until the best model is found with high confidence

$|C_i| = 14$

RANSAC for general model

A given model type has a *minimal set* – the smallest number of samples from which the model can be computed.

- Line: 2 points

RANSAC for general model

Image transformations are models. Minimal set of point pairs/matches:

- Translation: pick one pair of matched points
- Homography (for plane) – pick 4 point pairs
- Fundamental matrix – pick 8 point pairs (really 7 but lets not go there)

RANSAC for general model

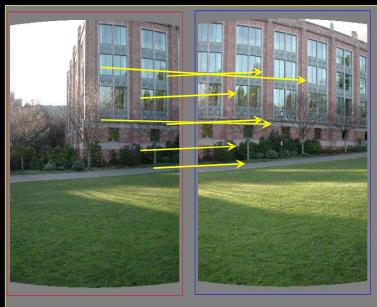
General RANSAC algorithm

- Randomly select S points (or point pairs) to form a *sample*
- Instantiate the model
- Get consensus set C_i : The points within error bounds (distance threshold) of the model
- If $|C_i| > T$, terminate and return model
- Repeat for N trials, return model with max $|C_i|$

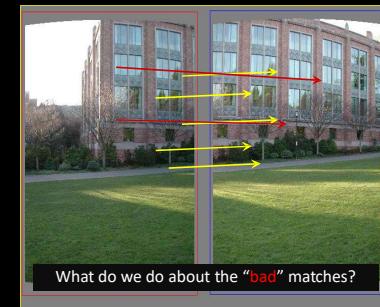
Choosing the parameters

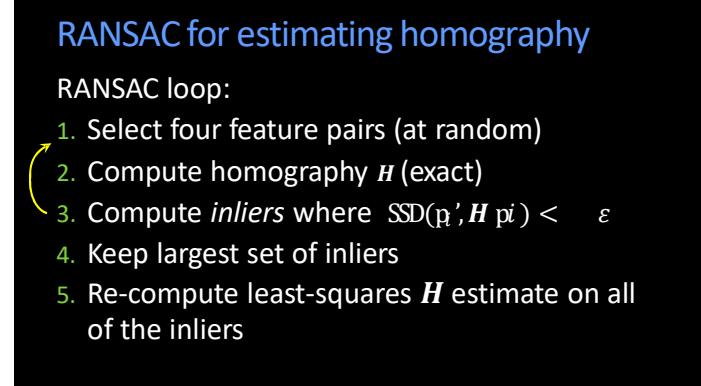
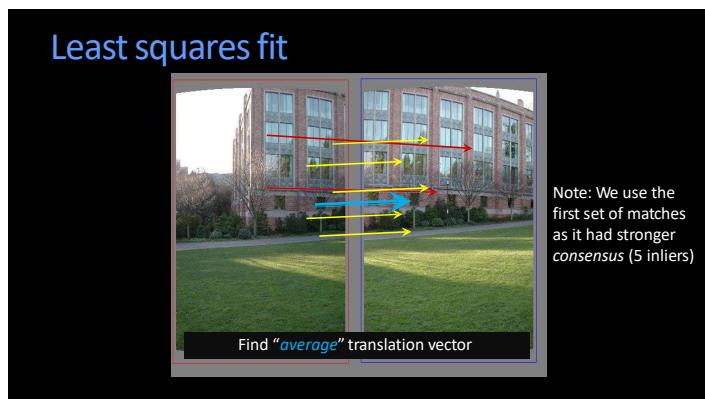
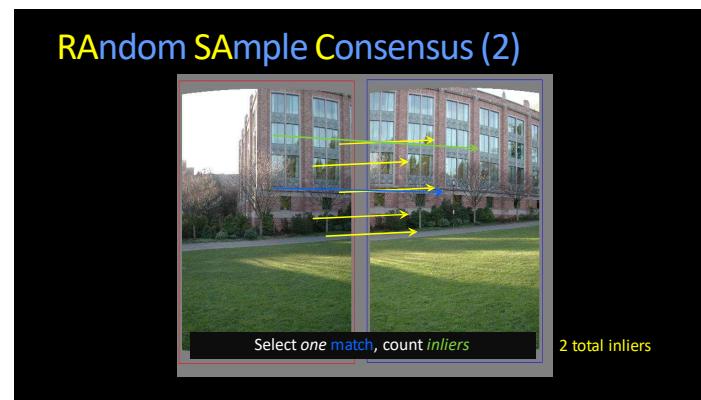
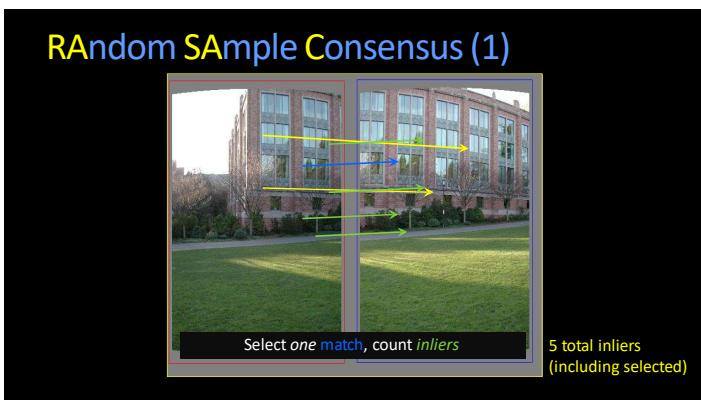
1. Initial number of points in the minimal set s
 - Typically minimum number needed to fit the model
2. Distance threshold t

Matching features



Matching features

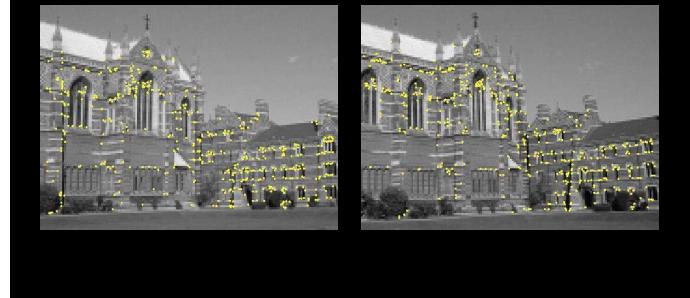




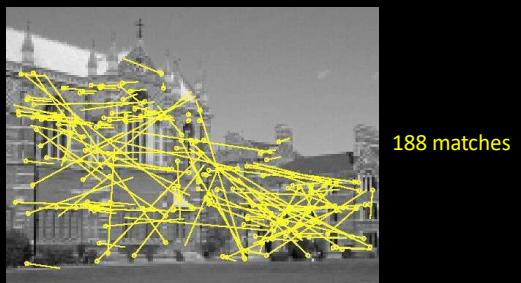
RANSAC for recognition



RANSAC for finding fundamental matrix

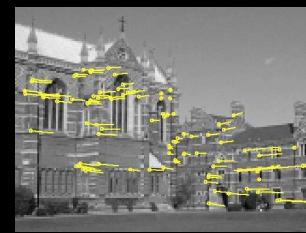


Putative matches (motion) by cross-correlation

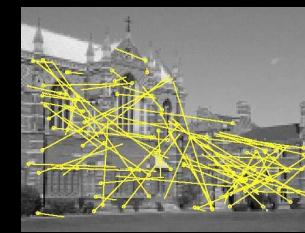


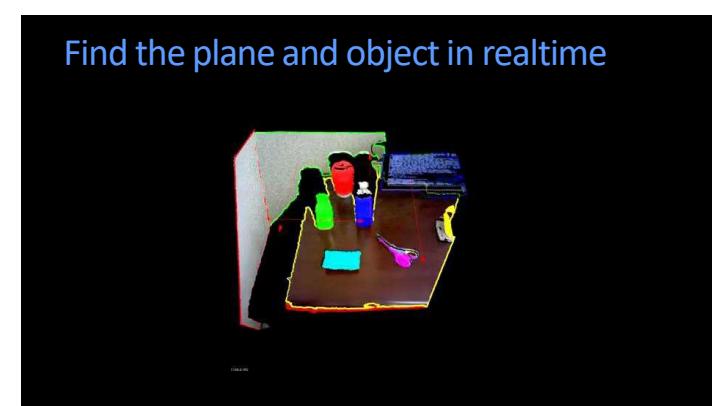
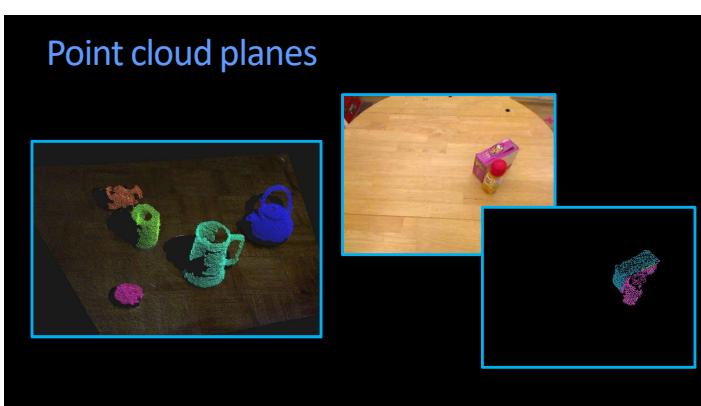
RANSAC for fundamental matrix

Inliers (99)



Outliers (89)





RANSAC: Conclusions

The good...

- Simple and general
- Applicable to many different problems, often works well in practice
- Robust to large numbers of outliers
- Applicable for larger number of parameters than Hough transform
- Parameters are easier to choose than Hough transform

RANSAC: Conclusions

The not-so-good...

- Computational time grows quickly with the number of model parameters
- Not as good for getting multiple fits
- Really not good for approximate models

RANSAC: Conclusions

Common applications

- Computing a homography (e.g., image stitching) or other image transform
- Estimating fundamental matrix (relating two views)
- *Pretty much every problem in robot vision*