

# Regularization

## L2

### Regularization

- Adding regularization to NN will help it reduce variance (overfitting)
- L1 matrix norm:
  - $\|w\| = \text{Sum}(|w[i,j]|)$  # sum of absolute values of all w
- L2 matrix norm because of arcane technical math reasons is called Frobenius norm:
  - $\|w\|^2 = \text{Sum}(|w[i,j]|^2)$  # sum of all w squared
  - Also can be calculated as  $\|w\|^2 = W.T * W$

## L2-Train

- Regularization for NN:
  - The normal cost function that we want to minimize is:
 
$$J(w_1, b_1, \dots, w_L, b_L) = (1/m) * \text{Sum}(L(y^{(1)}, y^{(1)}))$$
  - The L2 regularization version:
 
$$J(w, b) = (1/m) * \text{Sum}(L(y^{(1)}, y^{(1)})) + (\lambda/2m) * \text{Sum}(\|w^{(1)}\|^2)$$
  - We stack the matrix as one vector ( $m \times 1$ ) and then we apply  $\sqrt{w_1^2 + w_2^2 + \dots}$
  - To do back propagation (old way):
 
$$dw^{(1)} = (\text{from back propagation})$$
  - The new way:
 
$$dw^{(1)} = (\text{from back propagation}) + \lambda/m * w^{(1)}$$
  - So plugging it in weight update step:

```
w^{(1)} = w^{(1)} - learning\_rate * dw^{(1)}
= w^{(1)} - learning\_rate * ((from back propagation) + lambda/m * w^{(1)})
= w^{(1)} - (learning\_rate*lambda/m) * w^{(1)} - learning\_rate * (from back propagation)
= (1 - (learning\_rate*lambda/m)) * w^{(1)} - learning\_rate * (from back propagation)
```

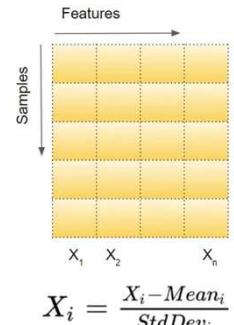
## Drop-Out

```
keep_prob = 0.8 # 0 <= keep_prob <= 1
l = 3 # this code is only for layer 3
# the generated number that are less than 0.8 will be dropped. 80% stay, 20% dropped
d3 = np.random.rand(a[1].shape[0], a[1].shape[1]) < keep_prob

a3 = np.multiply(a3, d3) # keep only the values in d3

# increase a3 to not reduce the expected value of output
# (ensures that the expected value of a3 remains the same) - to solve the scaling problem
a3 = a3 / keep_prob
```

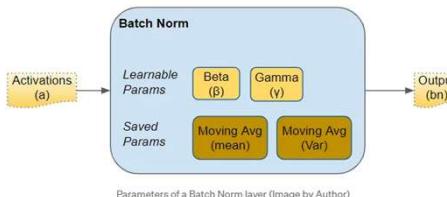
## Batch Normalization



How we normalize (Image by Author)

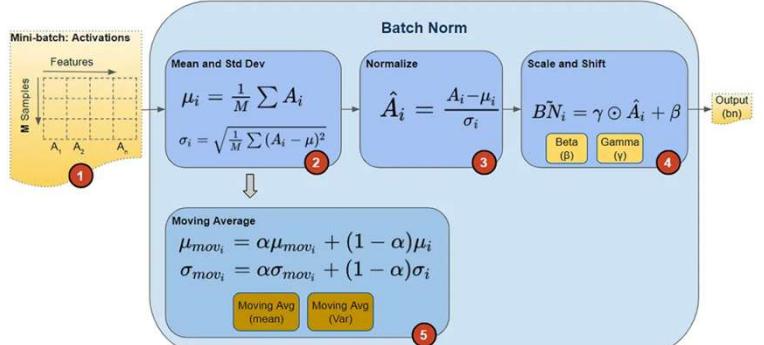
## Batch Normalization

- Two learnable parameters called beta and gamma.
- Two non-learnable parameters (Mean Moving Average and Variance Moving Average) are saved as part of the 'state' of the Batch Norm layer.



Parameters of a Batch Norm layer (Image by Author)

## Batch Normalization



Calculations performed by Batch Norm layer (Image by Author)

## Weighted Averages

```
v = 0
Repeat
{
    Get theta(t)
    v = beta * v + (1-beta) * theta(t)
}
```

## Momentum

### Gradient descent with momentum

- The momentum algorithm almost always works faster than standard gradient descent.
- The simple idea is to calculate the exponentially weighted averages for your gradients and then update your weights with the new values.
- Pseudo code:

```

vdl = 0, vdb = 0
on iteration t:
    # can be mini-batch or batch gradient descent
    compute dw, db on current mini-batch

    vdl = beta * vdl + (1 - beta) * dw
    vdb = beta * vdb + (1 - beta) * db
    W = W - learning_rate * vdl
    b = b - learning_rate * vdb
  
```

## RMS-Prop

### RMSprop

- Stands for Root mean square prop.
- This algorithm speeds up the gradient descent.
- Pseudo code:

```

sdW = 0, sdb = 0
on iteration t:
    # can be mini-batch or batch gradient descent
    compute dw, db on current mini-batch

    sdW = (beta * sdW) + (1 - beta) * dw^2 # squaring is element-wise
    sdb = (beta * sdb) + (1 - beta) * db^2 # squaring is element-wise
    W = W - learning_rate * dw / sqrt(sdW)
    b = b - learning_rate * db / sqrt(sdb)
  
```

## Adam optimization algorithm

- Stands for Adaptive Moment Estimation.
- Adam optimization and RMSprop are among the optimization algorithms that worked very well with a lot of NN architectures.
- Adam optimization simply puts RMSprop and momentum together!
- Pseudo code:

```

vdl = 0, vdb = 0
sdW = 0, sdb = 0
on iteration t:
    # can be mini-batch or batch gradient descent
    compute dw, db on current mini-batch

    vdl = (beta1 * vdl) + (1 - beta1) * dw # momentum
    vdb = (beta1 * vdb) + (1 - beta1) * db # momentum

    sdW = (beta2 * sdW) + (1 - beta2) * dw^2 # RMSprop
    sdb = (beta2 * sdb) + (1 - beta2) * db^2 # RMSprop

    vdl = vdl / (1 - beta1^t) # fixing bias
    vdb = vdb / (1 - beta1^t) # fixing bias

    sdW = sdW / (1 - beta2^t) # fixing bias
    sdb = sdb / (1 - beta2^t) # fixing bias

    W = W - learning_rate * vdl / (sqrt(sdW) + epsilon)
    b = b - learning_rate * vdb / (sqrt(sdb) + epsilon)
  
```

• Hyperparameters for Adam:

- Learning rate: needed to be tuned.
- `beta1`: parameter of the momentum - 0.9 is recommended by default.
- `beta2`: parameter of the RMSprop - 0.999 is recommended by default.
- `epsilon`:  $10^{-8}$  is recommended by default.

## AdamW

```

first_moment = 0
second_moment = 0
for t in range(1, num_iterations):
    dx = compute_gradient(x) # Standard Adam computes L2 here
    first_moment = beta1 * first_moment + (1 - beta1) * dx
    second_moment = beta2 * second_moment + (1 - beta2) * dx * dx
    first_unbias = first_moment / (1 - beta1 ** t)
    second_unbias = second_moment / (1 - beta2 ** t)
    x -= learning_rate * first_unbias / (np.sqrt(second_unbias) + 1e-7)) # AdamW (Weight Decay) adds term here
    x -= learning_rate * weight_decay * x
  
```

```
# 1 Compute the gradient
g_t = ∇L(W)

# 2 Update Adam's moment estimates
m_t = β1 * m_{t-1} + (1 - β1) * g_t
v_t = β2 * v_{t-1} + (1 - β2) * g_t**2

# 3 Apply bias correction
m_hat = m_t / (1 - β1**t)
v_hat = v_t / (1 - β2**t)

# 4 Update the weights (Adam update step)
W -= η * m_hat / (sqrt(v_hat) + ε)

# 5 Apply weight decay (separate from the gradient)
W -= η * λ * W
```